

# In-Class Project: ChIP-Seq Peak Calling with MACS3

Payal Priyadarshini

2025-05-05

## Introduction

Welcome! This project will guide you through the process of identifying potential protein binding sites or histone modifications in the genome using **simulated** ChIP-Seq data designed to have a strong signal. We will use a software called MACS3 (Model-based Analysis of ChIP-Seq version 3), a super useful tool for peak calling and several other types of analysis of ChIP-Seq data.

In this exercise, you will:

1. Generate simulated paired-end ChIP-Seq data with stronger enrichment signals.
2. Execute the **MACS3** peak caller on this simulated data from within an R script, attempting default model building.
3. Load the resulting `.narrowPeak` file into R using `rtracklayer` (if peaks were found).
4. Briefly inspect the `GRanges` object containing the peaks and visualize some basic properties (histograms).
5. Visualize the raw read coverage for treatment and control alongside the called peaks in a specific genomic region using `Gviz`.

**Dataset:** We will generate simulated BED files representing paired-end reads:

- Treatment: Simulates reads from a ChIP experiment with strong enrichment at specific locations.
- Control: Simulates background reads.
- Format: Paired-end reads in BED format.

**Paired-End Data:** Our simulated data represents paired-end sequencing. MACS3 will attempt to infer the fragment size distribution from the data itself.

**Prerequisite:** You **must** have MACS3 installed and accessible from your command line *before* proceeding. Test with `macs3 --version` in your terminal.

## 1. Setup R Environment

Simply run the code chunk below to install to ensure the necessary libraries for the rest of this script :) \* `dplyr`: Data manipulation. \* `rtracklayer`: Import BED files and MACS3 peak results \* `GenomicRanges`: Provides core infrastructure for genomic data \* `Gviz`: Required for plotting genomic tracks (coverage, peaks)

```
if (!requireNamespace("dplyr", quietly = TRUE)) install.packages("dplyr")
if (!requireNamespace("BiocManager", quietly = TRUE)) install.packages("BiocManager")
if (!requireNamespace("rtracklayer", quietly = TRUE)) BiocManager::install("rtracklayer")
# Gviz has many dependencies, installation might take a while
if (!requireNamespace("Gviz", quietly = TRUE)) BiocManager::install("Gviz")
# GenomicRanges is usually installed as a dependency, but good to be explicit
if (!requireNamespace("GenomicRanges", quietly = TRUE)) BiocManager::install("GenomicRanges")

# Load the libraries for this session
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(rtracklayer)

## Loading required package: GenomicRanges

## Loading required package: stats4

## Loading required package: BiocGenerics

##
## Attaching package: 'BiocGenerics'

## The following objects are masked from 'package:dplyr':
##
##   combine, intersect, setdiff, union

## The following objects are masked from 'package:stats':
##
##   IQR, mad, sd, var, xtabs

## The following objects are masked from 'package:base':
##
##   anyDuplicated, aperm, append, as.data.frame, basename, cbind,
##   colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,
##   get, grep, grepl, intersect, is.unsorted, lapply, Map, mapply,
##   match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,
##   Position, rank, rbind, Reduce, rownames, sapply, setdiff, sort,
##   table, tapply, union, unique, unsplit, which.max, which.min

## Loading required package: S4Vectors

##
## Attaching package: 'S4Vectors'

## The following objects are masked from 'package:dplyr':
##
##   first, rename

## The following object is masked from 'package:utils':
##
##   findMatches

```

```
## The following objects are masked from 'package:base':
##
##     expand.grid, I, unname

## Loading required package: IRanges

##
## Attaching package: 'IRanges'

## The following objects are masked from 'package:dplyr':
##
##     collapse, desc, slice

## Loading required package: GenomeInfoDb

## Warning: package 'GenomeInfoDb' was built under R version 4.3.3

library(GenomicRanges)
library(Gviz)

## Loading required package: grid
```

## 2. Simulate Paired-End ChIP-Seq Data

Normally, the `callpeaks` function of MACS3 would be used on `.bam` or `.bed` files. However, these files are quite large, so in order to save our computers from the torture of a gigantic download and prevent the need for dealing with the FU servers, we decided to simulate paired-end ChIP-Seq Data. If you are interested, there is information in the “Simulation\_Algo.txt” file in this project’s directory that explains the overall logic behind the simulation.

Run the code chunk below to generate the simulated ChIP-Seq data and answer the following questions:

1. If we were to use real ChIP-Seq data in this analysis, where could we normally find such data (e.g., which website/database)?
2. Which file formats are passable to the `callpeaks` function of MACS3?
3. What are the important data quality control measures one should take before attempting to call peaks?

```
# Sim params
set.seed(456)
n_pairs_treatment <- 8000
n_pairs_control <- 1500
read_length <- 50
fragment_mean_size <- 200
fragment_sd_size <- 25
genome_size <- 1e6
chr_name <- "chrSim"

# Helper function to generate reads
generate_read_pair <- function(pair_id, chr, center, frag_len, read_len, strand_choice = c("+", "-")) {
  frag_start <- max(0, round(center - frag_len / 2))
  frag_end <- frag_start + frag_len
```

```

if (sample(strand_choice, 1) == "+") {
  read1_start <- frag_start
  read1_end <- frag_start + read_len
  read2_start <- frag_end - read_len
  read2_end <- frag_end
  strand1 <- "+"
  strand2 <- "-"
} else {
  read1_start <- frag_end - read_len
  read1_end <- frag_end
  read2_start <- frag_start
  read2_end <- frag_start + read_len
  strand1 <- "-"; strand2 <- "+"
}
read1_start <- max(0, read1_start)
read1_end <- min(genome_size, read1_end)
read2_start <- max(0, read2_start)
read2_end <- min(genome_size, read2_end)
if (read1_start >= read1_end) read1_end <- read1_start + 1
if (read2_start >= read2_end) read2_end <- read2_start + 1
read1 <- data.frame(chr = chr, start = read1_start, end = read1_end,
                    name = paste0("sim_read_", pair_id, "/1"), score = 0, strand = strand1)
read2 <- data.frame(chr = chr, start = read2_start, end = read2_end,
                    name = paste0("sim_read_", pair_id, "/2"), score = 0, strand = strand2)
return(rbind(read1, read2))
}

# Generate treatment reads
hotspot_centers <- round(genome_size * c(0.15, 0.35, 0.65, 0.85))
n_hotspots <- length(hotspot_centers)
n_per_hotspot <- round(n_pairs_treatment * 0.90 / n_hotspots)
n_background <- n_pairs_treatment - (n_per_hotspot * n_hotspots)
hotspot_sd <- fragment_mean_size * 0.5

centers_treatment <- c(
  unlist(lapply(hotspot_centers, function(center) {
    rnorm(n_per_hotspot, mean = center, sd = hotspot_sd)
  })),
  runif(n_background, min = 0, max = genome_size)
)
centers_treatment <- pmax(0, pmin(genome_size, round(centers_treatment)))

frag_lengths_treatment <- round(rnorm(n_pairs_treatment, mean = fragment_mean_size, sd = fragment_sd_si
frag_lengths_treatment[frag_lengths_treatment < read_length * 2] <- read_length * 2

cat("Generating treatment reads...\n")

## Generating treatment reads...

treatment_reads_list <- lapply(1:n_pairs_treatment, function(i) {
  generate_read_pair(pair_id = i, chr = chr_name, center = centers_treatment[i],
                    frag_len = frag_lengths_treatment[i], read_len = read_length)
})

```

```
treatment_bed <- bind_rows(treatment_reads_list) %>% arrange(chr, start)
```

```
# Generate control reads
```

```
cat("Generating control reads...\n")
```

```
## Generating control reads...
```

```
centers_control <- round(runif(n_pairs_control, min = 0, max = genome_size))
```

```
centers_control <- pmax(0, pmin(genome_size, centers_control))
```

```
frag_lengths_control <- round(rnorm(n_pairs_control, mean = fragment_mean_size, sd = fragment_sd_size))
```

```
frag_lengths_control[frag_lengths_control < read_length * 2] <- read_length * 2
```

```
control_reads_list <- lapply(1:n_pairs_control, function(i) {  
  generate_read_pair(pair_id = i + n_pairs_treatment, chr = chr_name, center = centers_control[i],  
    frag_len = frag_lengths_control[i], read_len = read_length)  
})
```

```
control_bed <- bind_rows(control_reads_list) %>% arrange(chr, start)
```

```
# Write BED files
```

```
sim_input_dir <- "macs_simulated_input"
```

```
if (!dir.exists(sim_input_dir)) dir.create(sim_input_dir)
```

```
treatment_bed_file <- file.path(sim_input_dir, "treatment_strong.bed")
```

```
control_bed_file <- file.path(sim_input_dir, "control_strong.bed")
```

```
cat("Writing BED files...\n")
```

```
## Writing BED files...
```

```
write.table(treatment_bed, file = treatment_bed_file, sep = "\t", row.names = FALSE, col.names = FALSE,
```

```
write.table(control_bed, file = control_bed_file, sep = "\t", row.names = FALSE, col.names = FALSE, qu
```

```
cat("Simulated BED files created:\n")
```

```
## Simulated BED files created:
```

```
cat("- Treatment:", treatment_bed_file, "\n")
```

```
## - Treatment: macs_simulated_input/treatment_strong.bed
```

```
cat("- Control:", control_bed_file, "\n")
```

```
## - Control: macs_simulated_input/control_strong.bed
```

```
if (!file.exists(treatment_bed_file) || file.info(treatment_bed_file)$size == 0) stop("Failed to create
```

```
if (!file.exists(control_bed_file) || file.info(control_bed_file)$size == 0) stop("Failed to create or v
```

```
cat("BED files written successfully.\n")
```

```
## BED files written successfully.
```

### 3. Run MACS3 on Simulated Data

Now we construct and execute the `macs3 callpeak` command using the simulated BED files. We will let MACS3 attempt to build its model. Edit the code below where indicated.

```
# Define output directory for MACS3 results
macs3_output_dir <- "macs3_results_simulated"
if (!dir.exists(macs3_output_dir)) dir.create(macs3_output_dir)

# Define the output name prefix
output_name <- "simulated_run"

# TASK 3.1
# -----
# Construct the MACS3 command arguments for input into the system2() function
# for command line operations from R.
# Hints:
# 1. Appropriate genome size is "1e6".
# 2. Files containing simulated data are in .bed format.
# 3. The output directory for the MACS3 peak calling is specified in the variable macs3_output_dir.
# 4. The output name for the MACS3 peak calling output file is specified in the variable output_name.
# 5. A good default q-value is 0.05.
# 6. In our instance, we would like to bypass building a shifting model for
#    determining the shift of the reads towards the center.
macs3_args <- c(
  "callpeak",
  "-t", treatment_bed_file,
  "-c", control_bed_file,
  "-f", "BED",
  "-g", "1e6",
  "-n", chr_name,
  "--outdir", "/Users/payalpriyadarshini/Documents/Data_Science_2nd_sem/macs3_output_dir",
  "-q", "0.05",
  "--nomodel"
)

# TASK 3.2
# -----
# Execute MACS3:
# Edit the value of macs3_executable (below) to be the full file path to your MACS3 install.
# This can be found by calling "which macs3" in your Terminal.
macs3_executable <- "/opt/anaconda3/bin/macs3" # <-- EDIT THIS LINE with full path if needed
# -----

cat("Running MACS3 on simulated data...\n")
```

## Running MACS3 on simulated data...

```
full_command_string <- paste(macs3_executable, paste(macs3_args, collapse = " "))
```

```
# TASK 3.3
# -----
# Evoke a command line operation (using system2()) to call MACS3's callpeaks
```

```

# function.
# Hint: The macs3_executable variable is used as the first term in a MACS3 command line call
# Example: In "macs3 -f arg1 -n arg2 ..." macs3 (or its full path) is stored in macs3_executable above.

start_time <- Sys.time()
macs3_output <- system2(macs3_executable, args = macs3_args)      # YOUR CODE HERE
# -----
end_time <- Sys.time()
duration <- end_time - start_time

cat("\n----- MACS3 Output Log Start ----- \n")

##
## ----- MACS3 Output Log Start -----

#cat(paste(macs3_output, collapse = "\n"), "\n")
cat("----- MACS3 Output Log End ----- \n\n")

## ----- MACS3 Output Log End -----

cat("MACS3 run duration:", duration, "seconds\n\n")

## MACS3 run duration: 0.5642381 seconds

# Check if output files were created
expected_peak_file <- file.path(macs3_output_dir, paste0(output_name, "_peaks.narrowPeak"))
expected_xls_file <- file.path(macs3_output_dir, paste0(output_name, "_peaks.xls"))

```

## 4. Load and Explore MACS3 Peaks

Assuming MACS3 completed successfully and generated a `.narrowPeak` file, we load it into R. Please also answer the following questions:

1.

```

# Path to the narrowPeak file generated by MACS3
narrowPeak_file <- file.path("/Users/payalpriyadarshini/Documents/Data_Science_2nd_sem/macs3_output_dir",
peaks <- NULL # Initialize peaks object

# TASK 4
# -----
# Import the narrowPeak file (narrowPeak_file) generated by MACS3 using
# the rtracklayer library.

if (file.exists(narrowPeak_file)) {
  cat("Importing peaks from:", narrowPeak_file, "... \n")
  peaks <- import(narrowPeak_file) # YOUR CODE HERE
  if (!is.null(peaks)) {
    cat("\nSuccessfully imported", length(peaks), "peaks. \n")
  }
}

```

```

        cat("\nPeaks object summary (first 6):\n")
        show(head(peaks))
    } else {
        cat("\nPeak import failed despite file existing.\n")
    }
} else {
    cat("\nCould not find the MACS3 narrowPeak file to import:", narrowPeak_file, "\n")
    cat("Check the MACS3 log and the warning messages from the previous chunk.\n")
}

```

```

## Importing peaks from: /Users/payalpriyadarshini/Documents/Data_Science_2nd_sem/macs3_output_dir/chrS
##
## Successfully imported 4 peaks.
##
## Peaks object summary (first 6):
## GRanges object with 4 ranges and 6 metadata columns:
##      seqnames      ranges strand |      name      score signalValue
##      <Rle>      <IRanges> <Rle> | <character> <numeric> <numeric>
## [1] chrSim 149663-150387      * | chrSim_peak_1    5135    149.658
## [2] chrSim 349623-350420      * | chrSim_peak_2    4762    130.676
## [3] chrSim 649638-650357      * | chrSim_peak_3    4537    117.911
## [4] chrSim 849609-850378      * | chrSim_peak_4    5098    148.523
##      pValue      qValue      peak
##      <numeric> <numeric> <integer>
## [1] 519.510 513.511 348
## [2] 479.846 476.218 380
## [3] 457.159 453.798 363
## [4] 514.312 509.868 402
## -----
## seqinfo: 1 sequence from an unspecified genome; no seqlengths

```

## 5. Basic Peak Visualization (Optional Histograms)

This code visualizes the distribution of peak widths and scores. Simply run the code below to get the output, no coding required for you :)

```

if (!is.null(peaks) && length(peaks) > 0) {
    peak_widths <- width(peaks)
    reasonable_widths <- peak_widths[peak_widths < quantile(peak_widths, 0.99, na.rm = TRUE)]
    hist(reasonable_widths, breaks = 50, main = "Distribution of Peak Widths (up to 99th percentile)", xlab = "Peak Width")


    peak_scores <- mcols(peaks)$score
    hist(peak_scores, breaks = 50, main = "Distribution of Peak Scores", xlab = "Peak Score (-log10(qvalue))")

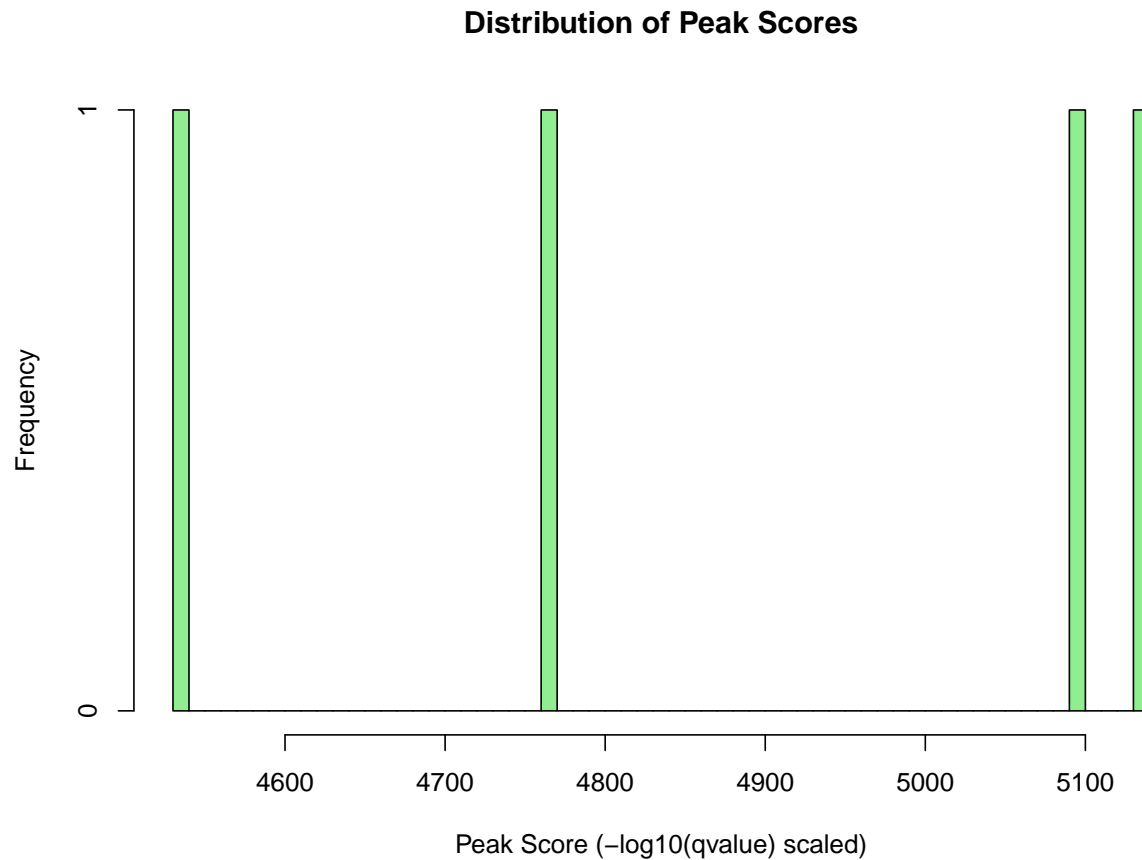
    cat("\nSummary statistics for loaded peaks:\n")
    cat("Widths:\n"); print(summary(peak_widths))
    cat("\nScores:\n"); print(summary(peak_scores))
} else {
    cat("\nNo peaks loaded, skipping histogram visualization.\n")
}

```



**Distribution of Peak Widths (up to 99th percentile)**





```
##
## Summary statistics for loaded peaks:
## Widths:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   720.0  723.8   747.5   753.2  777.0   798.0
##
## Scores:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   4537   4706   4930   4883   5107   5135
```

## 6. Visualize Read Coverage and Peaks

The code below generates a nice plot of where the peaks are located throughout the genome. Again, no coding needed as a reward for your hard work in the earlier tasks :) However, please answer the following questions:

1. Explain the difference in patterns between the control and treatment samples.
2. What would a control sample with significant peaks suggest about your prior ChIP-Seq experiment?

```
# Check if we have the necessary inputs for plotting
if (exists("treatment_bed_file") && file.exists(treatment_bed_file) &&
    exists("control_bed_file") && file.exists(control_bed_file) &&
```

```

exists("peaks") && !is.null(peaks) && length(peaks) > 0) {

cat("\nGenerating coverage plot...\n")

# Define region of interest
plot_center <- hotspot_centers[1]
plot_window <- 15000
plot_start <- max(0, plot_center - plot_window / 2)
plot_end <- min(genome_size, plot_center + plot_window / 2)
plot_chr <- chr_name

cat("Plotting region:", plot_chr, ":", plot_start, "-", plot_end, "\n")

# Import BED files as GRanges
cat("Importing BED files for coverage calculation...\n")
gr_treatment <- tryCatch(import(treatment_bed_file, format = "bed"), error=function(e) {message("Error im
gr_control <- tryCatch(import(control_bed_file, format = "bed"), error=function(e) {message("Error imp

if(is.null(gr_treatment) || is.null(gr_control)) {
  cat("Could not import BED files, skipping coverage plot.\n")
} else {
  # Calculate coverage
  cat("Calculating coverage...\n")
  cov_treatment <- coverage(gr_treatment)[[plot_chr]]
  cov_control <- coverage(gr_control)[[plot_chr]]

  # Get data for ylim calculation
  # Ensure region indices are valid before subsetting coverage vectors
  max_coord_treat <- length(cov_treatment)
  max_coord_ctrl <- length(cov_control)
  valid_indices_treat <- seq(max(1, plot_start), min(max_coord_treat, plot_end))
  valid_indices_ctrl <- seq(max(1, plot_start), min(max_coord_ctrl, plot_end))

  # Handle cases where the plot region might be outside the actual data range
  if(length(valid_indices_treat) == 0 || length(valid_indices_ctrl) == 0) {
    cat("Plotting region is outside the range of the coverage data. Skipping plot.\n")
  } else {
    # Extract numeric coverage data for the plot region
    plot_data_treat <- as.numeric(cov_treatment[valid_indices_treat])
    plot_data_ctrl <- as.numeric(cov_control[valid_indices_ctrl])

    # Calculate max coverage for ylim *from the numeric vectors*
    max_cov <- max(max(plot_data_treat), max(plot_data_ctrl), 1) # Ensure ylim >= 1

    # --- Create Gviz Tracks ---
    gtrack <- GenomeAxisTrack()

    dtrack_treatment <- DataTrack(start = valid_indices_treat, width = 1,
                                  chromosome = plot_chr, genome = "simG",
                                  data = plot_data_treat, # Use pre-calculated vector
                                  type = "histogram", name = "Treatment Cov",
                                  background.title = "darkblue", col.histogram = "darkblue", fill

```

```

dtrack_control <- DataTrack(start = valid_indices_ctrl, width = 1,
                             chromosome = plot_chr, genome = "simG",
                             data = plot_data_ctrl, # Use pre-calculated vector
                             type = "histogram", name = "Control Cov",
                             background.title = "darkred", col.histogram = "darkred", fill.his


atrack_peaks <- AnnotationTrack(peaks, name = "MACS3 Peaks",
                                chromosome = plot_chr,
                                background.title = "darkgreen", fill = "darkgreen", col="darkg
                                shape="box")

# --- Plot Tracks ---
cat("Plotting tracks...\n")
plotTracks(list(gtrack, dtrack_treatment, dtrack_control, atrack_peaks),
            from = plot_start, to = plot_end, chromosome = plot_chr,
            main = paste("Read Coverage and MACS3 Peaks around Hotspot 1"),
            ylim = c(0, max_cov * 1.1), # Use pre-calculated max_cov
            cex.main = 1
            )
cat("Coverage plot generated.\n")
} # End check for valid indices
} # End if BED files imported
} else {
cat("\nSkipping coverage plot because peaks were not loaded or BED files are missing.\n")
if (!exists("peaks")) cat(" -> Reason: 'peaks' object does not exist.\n")
else if (is.null(peaks)) cat(" -> Reason: 'peaks' object is NULL.\n")
else if (length(peaks) == 0) cat(" -> Reason: 'peaks' object is empty (length is 0).\n")
if (!exists("treatment_bed_file") || !file.exists(treatment_bed_file)) cat(" -> Reason: Treatment BED
if (!exists("control_bed_file") || !file.exists(control_bed_file)) cat(" -> Reason: Control BED file

}

##
## Generating coverage plot...
## Plotting region: chrSim : 142500 - 157500
## Importing BED files for coverage calculation...
## Calculating coverage...
## Plotting tracks...

```



```
## Coverage plot generated.
```

## Conclusion

In this exercise, you have: 1. Generated simulated paired-end ChIP-Seq data with stronger enrichment signals. 2. Executed the **MACS3** peak caller on this simulated data from within an R script, attempting default model building. 3. Loaded the resulting `.narrowPeak` file into R using `rtracklayer` (if peaks were found). 4. Briefly inspected the `GRanges` object containing the peaks and visualized some basic properties (histograms). 5. Visualized the raw read coverage for treatment and control alongside the called peaks in a specific genomic region using `Gviz`. ““